

**Network Computing & Programming**

**Lecture 2**

**Network Programming**

**Fundamentals**

**2019-2020-1**


**Slides are modified from Xiang Zhang**

# Content


- ❑ Concurrency & Parallelism
- ❑ Byte Ordering (Endianness)
- ❑ ILP/LP/LLP Definition
- ❑ Memory Alignment
- ❑ Protocol Header Data Structure
- ❑ Development & Debug Tools

# Concurrency & Parallelism

- ❑ **Parallelism** (并行)
- ❑ **Concurrency** (并发, 看起来同时发生)



并行模型



并发模型

# Concurrency in Networks & Hosts

- ❑ Peer-to-peer processes between two hosts work concurrently.
- ❑ All clients in the same host work concurrently.
- ❑ All requests introduced into a server can be processed concurrently.
- ❑ .....

# Concurrence & Parallelism

## ❑ Multiprogramming

- ❖ Tasks multiplex their executions on a single processor

## ❑ Multiprocessing

- ❖ Tasks multiplex their executions on a multiprocessor system where there is access to shared memory

## ❑ Distributed Processing

- ❖ Tasks multiplex their executions on several processors which do not share memory

# Byte Ordering

- Different computer architectures use different byte ordering to represent multi-byte values.

## Little-Endian



Addr A    Addr A+1

IBM 80x86

DEC VAX

DEC PDP-11

## Big-Endian



Addr A    Addr A+1

IBM 370

Motorola 68000

Sun

# Byte Ordering

- Suppose a Big Endian machine sends a 16 bit integer with the value 2:

0000000000000010

- A Little Endian machine will think it got the number 512:

0000001000000000

# Network Byte Order

- How do lower level layers communicate if they all represent values differently ?  
(data length fields in headers)
- A fixed byte order is used (called *network byte order*) for all control data.
  - ❖ TCP/IP : big-endian order



# Byte Order Dection

```
#include "stdio.h"
int main(int argc, char **argv)
{
    union{
        short    s;
        char     c[sizeof(short)];
    } un;
    un.s = 0x0102;
    if (sizeof(short) == 2) {
        if (un.c[0] == 1 && un.c[1] == 2)
            printf("big-endian\n");
        else if (un.c[0] == 2 && un.c[1] == 1)
            printf("little-endian\n");
        else printf("unknown\n");
    }
    else
        printf("sizeof(short) = %d\n",
sizeof(short));
    exit(0);
}
```

# Byte Order Conversion

■ Convert multi-byte integer types from host byte order to network byte order

- ① htons()          host to network short
- ② htonl()          host to network long
- ③ ntohs()          network to host short
- ④ ntohl()          network to host long

■ <arpa/inet.h>

# Quiz

- ❑ Please write out 192.168.128.64 by Big-Endian & Little-Endian & Network Byte Order rules separately.
- ❑ If we want to send out a message with a destination IP address 222.111.0.1, which is stored by Little-Endian rule, what shall we do?

# Quiz

- ❑ Two struct definitions
  - ❑ `struct { char a; char b; };`
  - ❑ `struct { int a; int b; };`
- ❑ Considering Big-Endian and Little-Endian rules separately.
  - ❑ If to read memory from the low address to the high address by the size of type, is a available firstly or secondly?
  - ❑ If to transmit bytes into network from the low address to the high address one by one, is a transmitted firstly or secondly?

# Data Type Models: ILP/LP/LLP

## □ ILP

- ❖ integers (I), long integers (L), and pointers (P)
- ❖ Linux/Unix: LP64; Windows: LLP64

	LP32	ILP32	LP64	LLP64	ILP64
char	8	8	8	8	8
short	16	16	16	16	16
int	16	32	32	32	64
long	32	32	64	32	64
long long	64	64	64	64	64
pointer	32	32	64	64	64

# Memory Alignment

## □ Take a look at the codes

```
#include <stdio.h>
```

```
struct Test
```

```
{
```

```
    char x1;
```

```
    int x2;
```

```
};
```

```
int main(int argc, char *argv[])
```

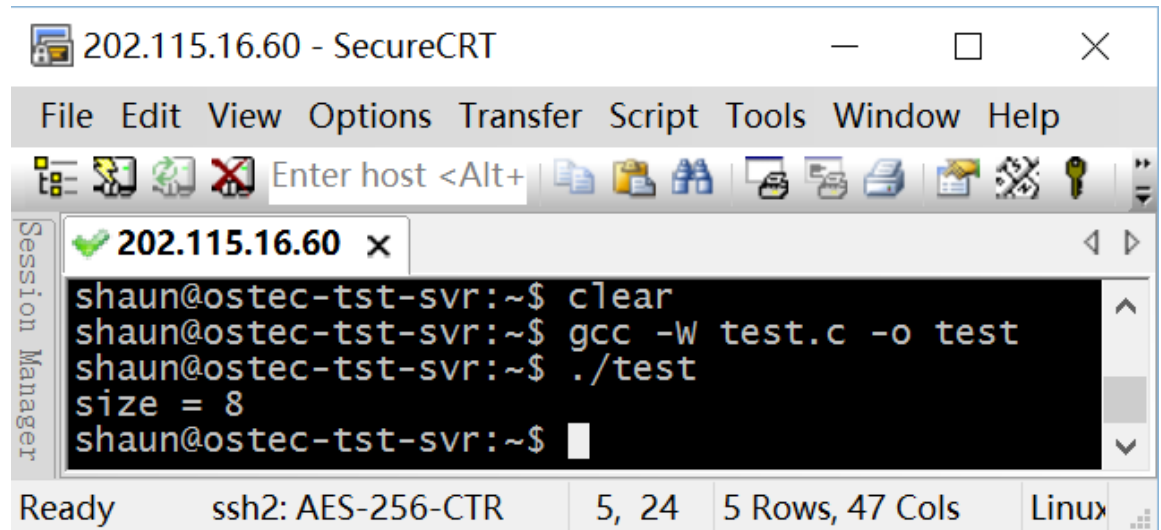
```
{
```

```
    int size = sizeof(struct Test);
```

```
    printf("size = %d\n", size);
```

```
    return 0;
```

```
}
```



202.115.16.60 - SecureCRT

File Edit View Options Transfer Script Tools Window Help

Enter host <Alt+>

Session Manager

202.115.16.60 x

```
shaun@ostec-tst-svr:~$ clear
shaun@ostec-tst-svr:~$ gcc -W test.c -o test
shaun@ostec-tst-svr:~$ ./test
size = 8
shaun@ostec-tst-svr:~$
```

Ready ssh2: AES-256-CTR 5, 24 5 Rows, 47 Cols Linux

# Memory Alignment

- ❑ How to read a 4-byte value stored at an address that is not a multiple of 4 bytes?
  - ❖ i386 permits this kind of access. But NOT ALL architectures permit it
    - Can raise exceptions
- ❑ Portability of data structures
  - ❖ Compiler rearranges structure fields to be aligned according to platform-specific conventions
  - ❖ Automatically add padding to make things aligned
    - May no longer match the intended format

# Memory Alignment

- Another example, consider the following structure on a 32-bit machine

```
struct animal_struct
{
    char dog;           /* 1 byte    */
    unsigned long cat;  /* 4 bytes   */
    unsigned short pig; /* 2 bytes   */
    char fox;           /* 1 byte    */
};
```



# Memory Alignment

- Structure not laid out like that in memory
  - Natural alignment of structure's members is inefficient
- Instead, compiler creates padding

```
struct animal_struct {  
    char dog;                /* 1 byte */  
    u8 __pad0[3];            /* 3 bytes */  
    unsigned long cat;        /* 4 bytes */  
    unsigned short pig;       /* 2 bytes */  
    char fox;                 /* 1 byte */  
    u8 __pad1;                /* 1 byte */  
};
```

# Memory Alignment

- You can often rearrange the order of members in a structure to obviate the need for padding

```
struct animal_struct {  
    unsigned long cat; /* 4 bytes */  
    unsigned short pig; /* 2 bytes */  
    char dog;          /* 1 byte  */  
    char fox;          /* 1 byte  */  
};
```

# Memory Alignment


- ❑ Another option is to tell the compiler to pack the data structure with **NO PADDING** added
- ❑ Example: `<linux/edd.h>`

```
struct {  
    u16 id;  
    u64 lun;  
    u16 reserved1;  
    u32 reserved2;  
} __attribute__((packed)) scsi;
```


- ❑ Quiz: What is the size of `scsi` on systems with ILP32 or LP64?

# Memory Alignment

- ❑ No compiler optimizations



- ❑ `__attribute__((packed))`



- ❑ Some compiler optimizations



# Memory Alignment

## □ **#pragma pack**

- ❖ Each system platform has its own Memory Alignment Coefficient.
- ❖ Preprocessor directives can help to alter this value.

• **#pragma pack(n), n=1,2,4,8,16 bytes.**

## □ Rule:

- We need to carefully design the data structure to make full use of the memory and try not padding when doing network programming.

# Quiz

❑ What are the exact size on x32 Linux system with default pack number 4.

- `struct X1 { char a; int b; char c; };`
- `struct X2 { int a; char b; };`
- `struct X3 { char a; short b; };`
- `struct X4 { char a; short b; char c; };`
- `struct X5 { char a; long b; };`
- `struct X6 { char a; long b; char c; };`
- `struct X7 { char a; long long b; };`
- `struct X8 { char a; long long b; char c; };`
- `struct X9 { char a; int b; short c; };`
- `struct X10 { char a; short b; char c; int d; };`

# Protocol Header Data Structure-Ethernet



<linux/if\_ether.h>

```
struct ethhdr {  
    unsigned char    h_dest[ETH_ALEN];  
        /* destination eth addr */  
    unsigned char    h_source[ETH_ALEN];  
        /* source ether addr    */  
    __be16           h_proto;  
        /* packet type ID field */  
} __attribute__((packed));
```

# Protocol Header Data Structure-IP

1 byte

1 byte

1 byte

1 byte

VERS	HL	Service	Fragment Length	
Datagram ID			FLAG	Fragment Offset
TTL		Protocol	Header Checksum	
Source Address				
Destination Address				
Options (if any)				
Data				



# Protocol Header Data

## Structure-IP

<linux/ip.h>

```
struct iphdr {  
    #if defined(__LITTLE_ENDIAN_BITFIELD)  
        __u8 ihl:4, version:4;  
    #elif defined (__BIG_ENDIAN_BITFIELD)  
        __u8 version:4, ihl:4;  
    #else  
        #error "Please fix <asm/byteorder.h>"  
    #endif  
  
    __u8  tos;                /*服务类型*/  
    __be16 tot_len;           /*总长度*/  
    __be16 id;                /*标识*/  
    __be16 frag_off;          /*片偏移*/  
    __u8  ttl;                /*生存时间*/  
    __u8  protocol;           /*协议类型*/  
    __u16 check;              /*头部校验和*/  
    __be32 saddr;             /*源IP地址*/  
    __be32 daddr;             /*目的IP地址*/  
  
};
```

# Protocol Header Data Structure-TCP

1 byte

1 byte

1 byte

1 byte

Source Port			Destination Port		
Sequence Number					
Request Number					
offset	Reser.	Control		Window	
Checksum			Urgent Pointer		
Options (if any)					
Data					

# Protocol Header Data <linux/tcp.h>

## Structure-TCP

```
struct tcphdr {
    __be16 source;
    __be16 dest;
    __be32 seq;
    __be32 ack_seq;
#if
defined(__LITTLE_ENDIAN_BITFIELD)
    __u16    res1:4,
            doff:4,
            fin:1,
            syn:1,
            rst:1,
            psh:1,
            ack:1,
            urg:1,
            ece:1,
            cwr:1;
    __u16    doff:4,
            res1:4,
            cwr:1,
            ece:1,
            urg:1,
            ack:1,
            psh:1,
            rst:1,
            syn:1,
            fin:1;
#else
#error "Adjust your
<asm/byteorder.h> defines"
#endif
    __be16 window;
    __be16 check;
    __be16 urg_ptr;
};
```

# Protocol Header Data

## Structure-UDP

	1 byte	1 byte	1 byte	1 byte
	Source Port		Destination Port	
	Length		Checksum	
<linux/udp.h>	Data			

```
struct udphdr {
    __be16  source;
    __be16  dest;
    __be16  len;
    __sum16 check;
};
```

# Development & Debug Tools

- ❑ VMWare / Virtual box + Ubuntu Server
- ❑ GCC + GDB / Clang + LLDB
- ❑ netstat, ping, tcpdump

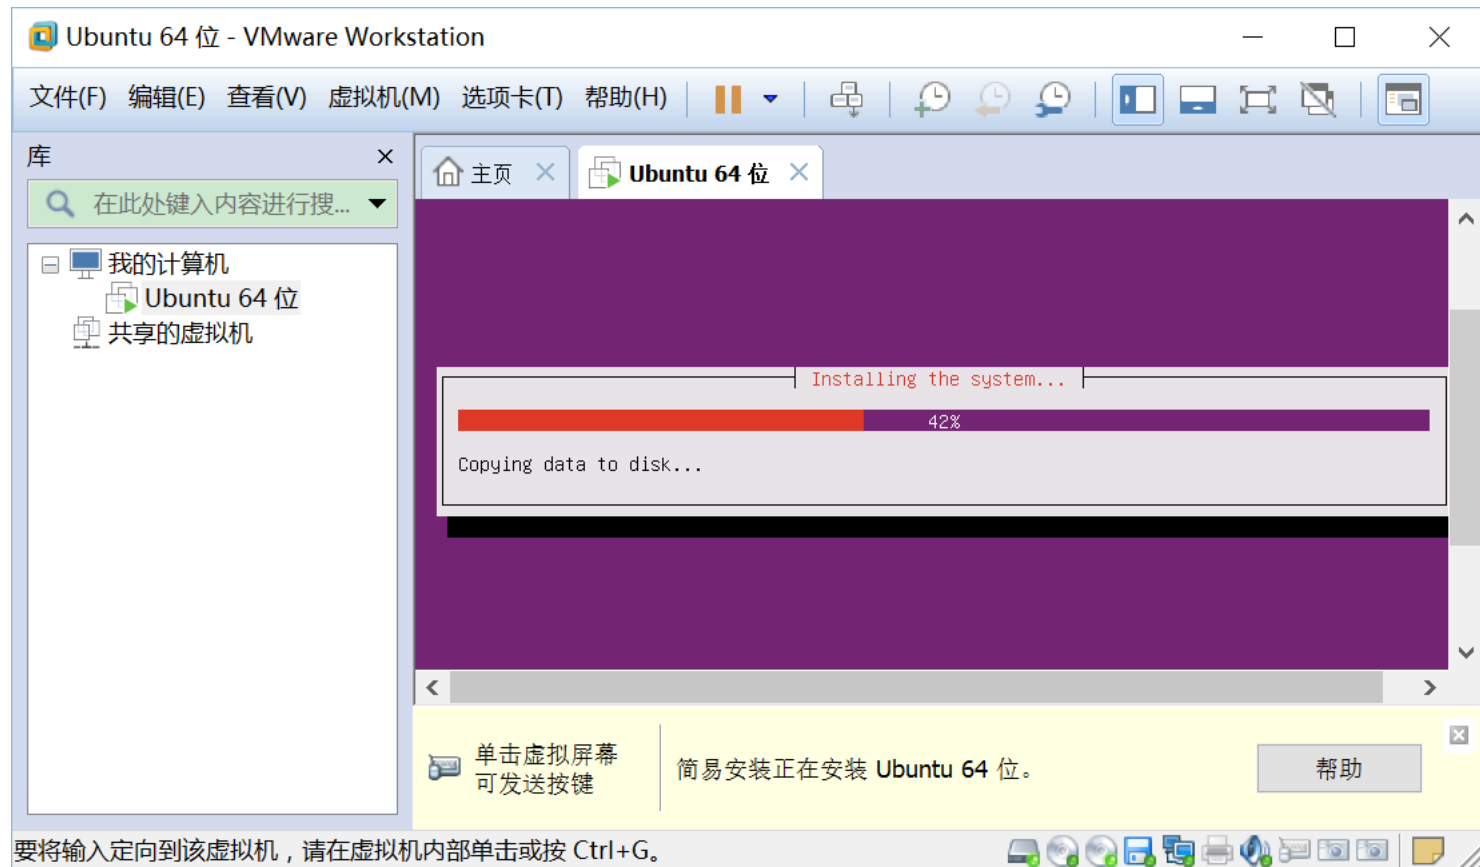
# Development & Debug Tools

- ❑ VMware Workstation
- ❑ VMware Fusion
- ❑ Virtual Box
- ❑ Ubuntu Server



# Development & Debug Tools

## ❑ Ubuntu x64



# Development & Debug Tools

## □ **GCC (GNU Compiler Collection)**


- ❖ Includes front ends for C, C++, Objective-C, Fortran, Ada, and Go, as well as libraries for these languages (libstdc++,...).

- ❖ <http://gcc.gnu.org>

## □ **LLVM (Low Level Virtual Machine)**

- ❖ Written in C++ and is designed for compile-time, link-time, run-time, and "idle-time" optimization of programs written in arbitrary programming languages

- ❖ Clang: an "LLVM native" C/C++/Objective-





# Development & Debug Tools

## □ **GDB (The GNU Project Debugger)**

- ❖ The program being debugged can be written in Ada, C, C++, Objective-C, Pascal (and many other languages).
- ❖ <http://www.gnu.org/software/gdb/>

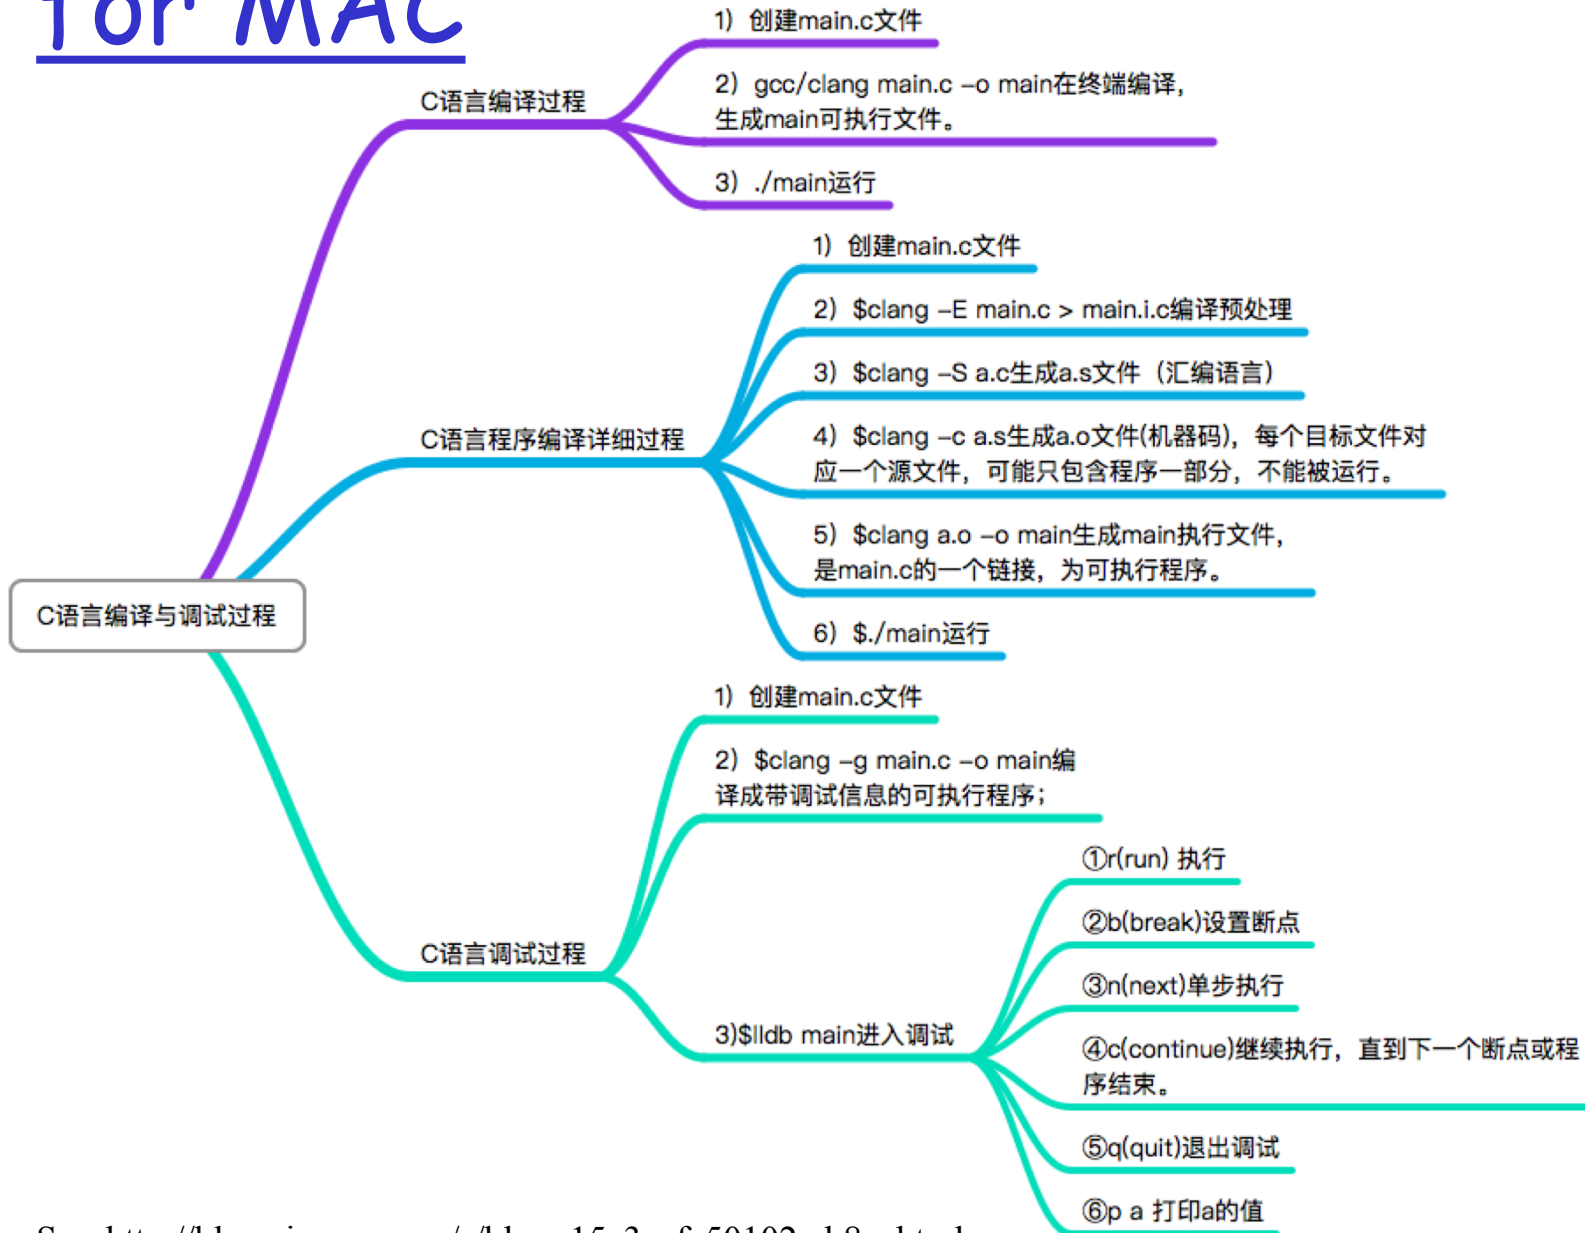
## □ **LLDB (Low Level Debugger)**

- ❖ LLDB project builds on libraries provided by LLVM and Clang to provide a great native debugger.
- ❖ It is also blazing fast and much more memory efficient than GDB at loading symbols.
- ❖ <http://lldb.llvm.org>

## □ **GDB TO LLDB COMMAND MAP**

- ❖ <http://lldb.llvm.org/lldb-gdb.html>

# for MAC



# Development & Debug Tools

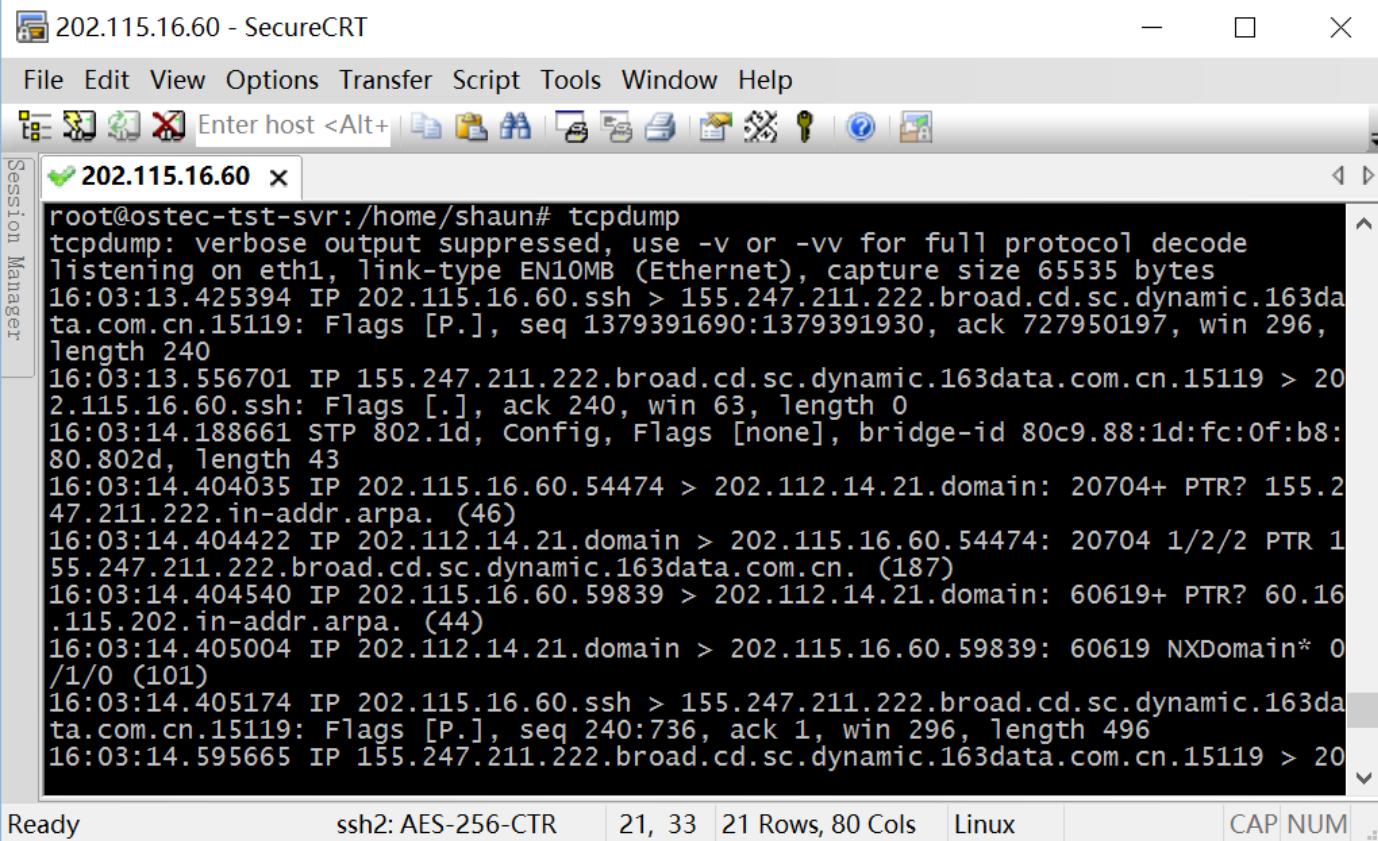
## □ netstat

```
root@ostec-tst-svr:/home/shaun# netstat -an
Active Internet connections (servers and established)
Proto Recv-Q Send-Q Local Address           Foreign Address         State
tcp        0      0 127.0.0.1:3306          0.0.0.0:*               LISTEN
tcp        0      0 0.0.0.0:22              0.0.0.0:*               LISTEN
tcp        0      0 202.115.16.60:22        115.28.233.246:56706    ESTABLISHED
tcp        0      416 202.115.16.60:22        222.211.247.155:15119  ESTABLISHED
tcp6       0      0 :::22                   :::*                    LISTEN
tcp6       0      0 :::8888                  :::*                    LISTEN
udp        0      0 0.0.0.0:11839           0.0.0.0:*               LISTEN
udp        0      0 0.0.0.0:68              0.0.0.0:*               LISTEN
udp6       0      0 :::61119                 :::*                    LISTEN

Active UNIX domain sockets (servers and established)
Proto RefCnt Flags       Type       State      I-Node    Path
unix   9      [ ]        DGRAM                    10500     /dev/log
unix   2      [ ACC ]    STREAM    LISTENING   11549     /var/run/dbus/system_
bus_socket
unix   2      [ ACC ]    STREAM    LISTENING   11840     /var/run/acpid.socket
unix   2      [ ACC ]    STREAM    LISTENING   9097      @/com/ubuntu/upstart
unix   2      [ ACC ]    STREAM    LISTENING   32910     /var/run/mysqld/mysq
d.sock
unix   2      [ ACC ]    SEQPACKET LISTENING   9193      /run/udev/control
unix   3      [ ]        STREAM    CONNECTED   2039
unix   3      [ ]        STREAM    CONNECTED   10481
```

# Development & Debug Tools

## ❑ tcpdump



The screenshot shows a SecureCRT terminal window titled "202.115.16.60 - SecureCRT". The terminal displays the output of the `tcpdump` command run on a host named `ostec-tst-svr`. The output shows network traffic on the `eth1` interface, including SSH connections and STP (Spanning Tree Protocol) frames. The status bar at the bottom indicates the session is "Ready", the encryption is "ssh2: AES-256-CTR", and the window size is "21, 33" with "21 Rows, 80 Cols".

```
root@ostec-tst-svr:/home/shaun# tcpdump
tcpdump: verbose output suppressed, use -v or -vv for full protocol decode
listening on eth1, link-type EN10MB (Ethernet), capture size 65535 bytes
16:03:13.425394 IP 202.115.16.60.ssh > 155.247.211.222.broad.cd.sc.dynamic.163da
ta.com.cn.15119: Flags [P.], seq 1379391690:1379391930, ack 727950197, win 296,
length 240
16:03:13.556701 IP 155.247.211.222.broad.cd.sc.dynamic.163data.com.cn.15119 > 20
2.115.16.60.ssh: Flags [.], ack 240, win 63, length 0
16:03:14.188661 STP 802.1d, Config, Flags [none], bridge-id 80c9.88:1d:fc:0f:b8:
80.802d, length 43
16:03:14.404035 IP 202.115.16.60.54474 > 202.112.14.21.domain: 20704+ PTR? 155.2
47.211.222.in-addr.arpa. (46)
16:03:14.404422 IP 202.112.14.21.domain > 202.115.16.60.54474: 20704 1/2/2 PTR 1
55.247.211.222.broad.cd.sc.dynamic.163data.com.cn. (187)
16:03:14.404540 IP 202.115.16.60.59839 > 202.112.14.21.domain: 60619+ PTR? 60.16
.115.202.in-addr.arpa. (44)
16:03:14.405004 IP 202.112.14.21.domain > 202.115.16.60.59839: 60619 NXDomain* 0
/1/0 (101)
16:03:14.405174 IP 202.115.16.60.ssh > 155.247.211.222.broad.cd.sc.dynamic.163da
ta.com.cn.15119: Flags [P.], seq 240:736, ack 1, win 296, length 496
16:03:14.595665 IP 155.247.211.222.broad.cd.sc.dynamic.163data.com.cn.15119 > 20
```