

**Network Computing & Programming**

**Lecture 3**

**Socket Programming**

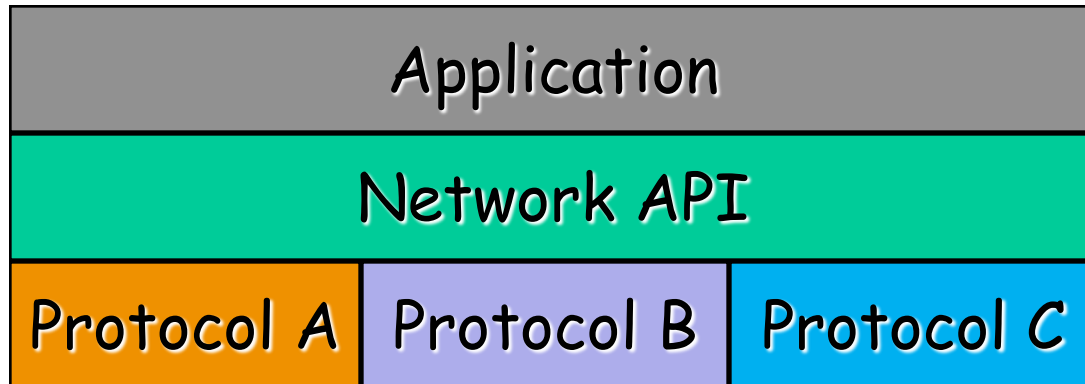
**– First Stage**

**2019-2020-1**

**Slides are modified from Xiang(Shaun) Zhang**

# Network Application Programming Interface (API)

- ❑ Services that provide the interface between application and protocol software (often by the operating system).



# Network API wish list

## ❑ **Generic Programming Interface**

- ❖ Support multiple communication protocol suites (families).
- ❖ Address (endpoint) representation independence.
- ❖ Provide special services for Client and Server.
- ❖ Support for message oriented and connection oriented communication.
- ❖ Work with existing I/O services (when this makes sense).
- ❖ Operating System independence.

# TCP/IP

- ❑ TCP/IP does not include an API definition.
- ❑ There are a variety of APIs for use with TCP/IP:
  - ❖ **Berkeley sockets**
    - API for Internet sockets and Unix domain sockets (used for inter-process communication (IPC))
  - ❖ XTI(X/Open Transport Interface)/**TLI** (Transport Layer Interface), Unix System V
  - ❖ **Winsock**, Windows
  - ❖ **MacTCP**, Mac
  - ❖ **Linux/POSIX Socket**, basically Berkeley sockets
    - Portable Operating System Interface, POSIX

# Functions needed:

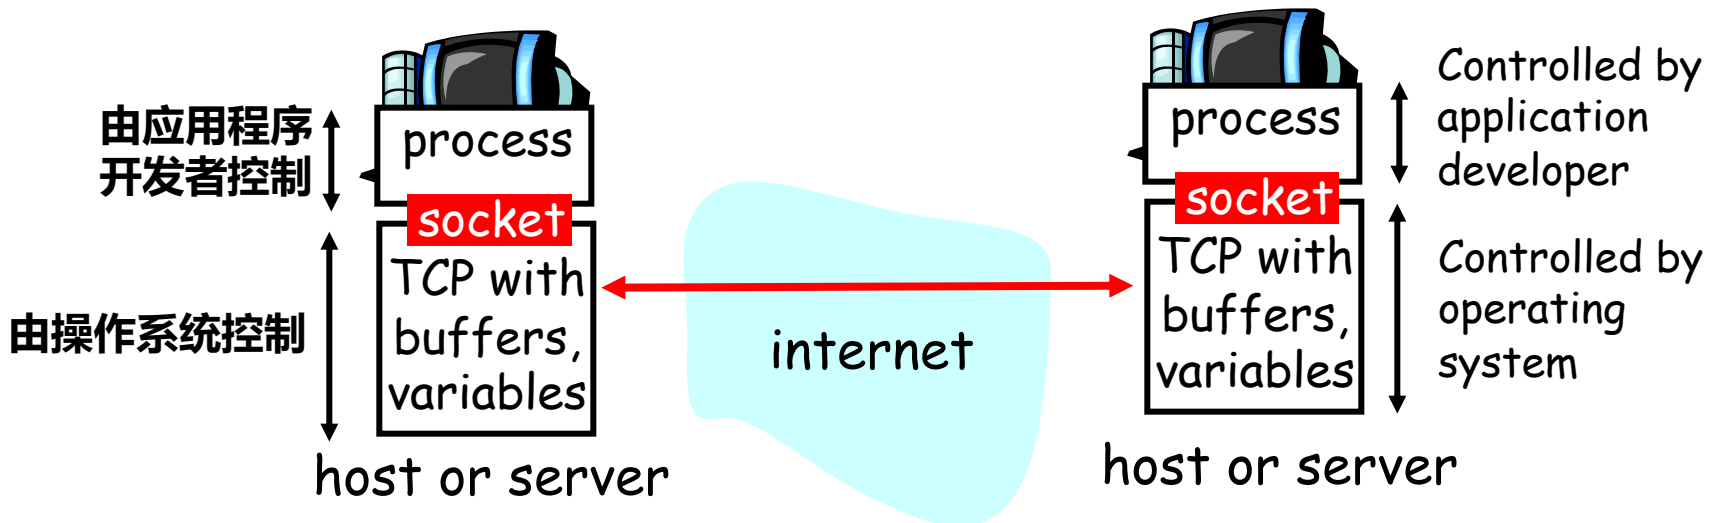
- ❑ Specify local and remote communication endpoints
- ❑ Initiate a connection
- ❑ Wait for incoming connection
- ❑ Send and receive data
- ❑ Terminate a connection gracefully
- ❑ Error handling

# Berkeley Sockets

## □ Generic:

- ❖ Support for multiple protocol families.
- ❖ Address representation independence

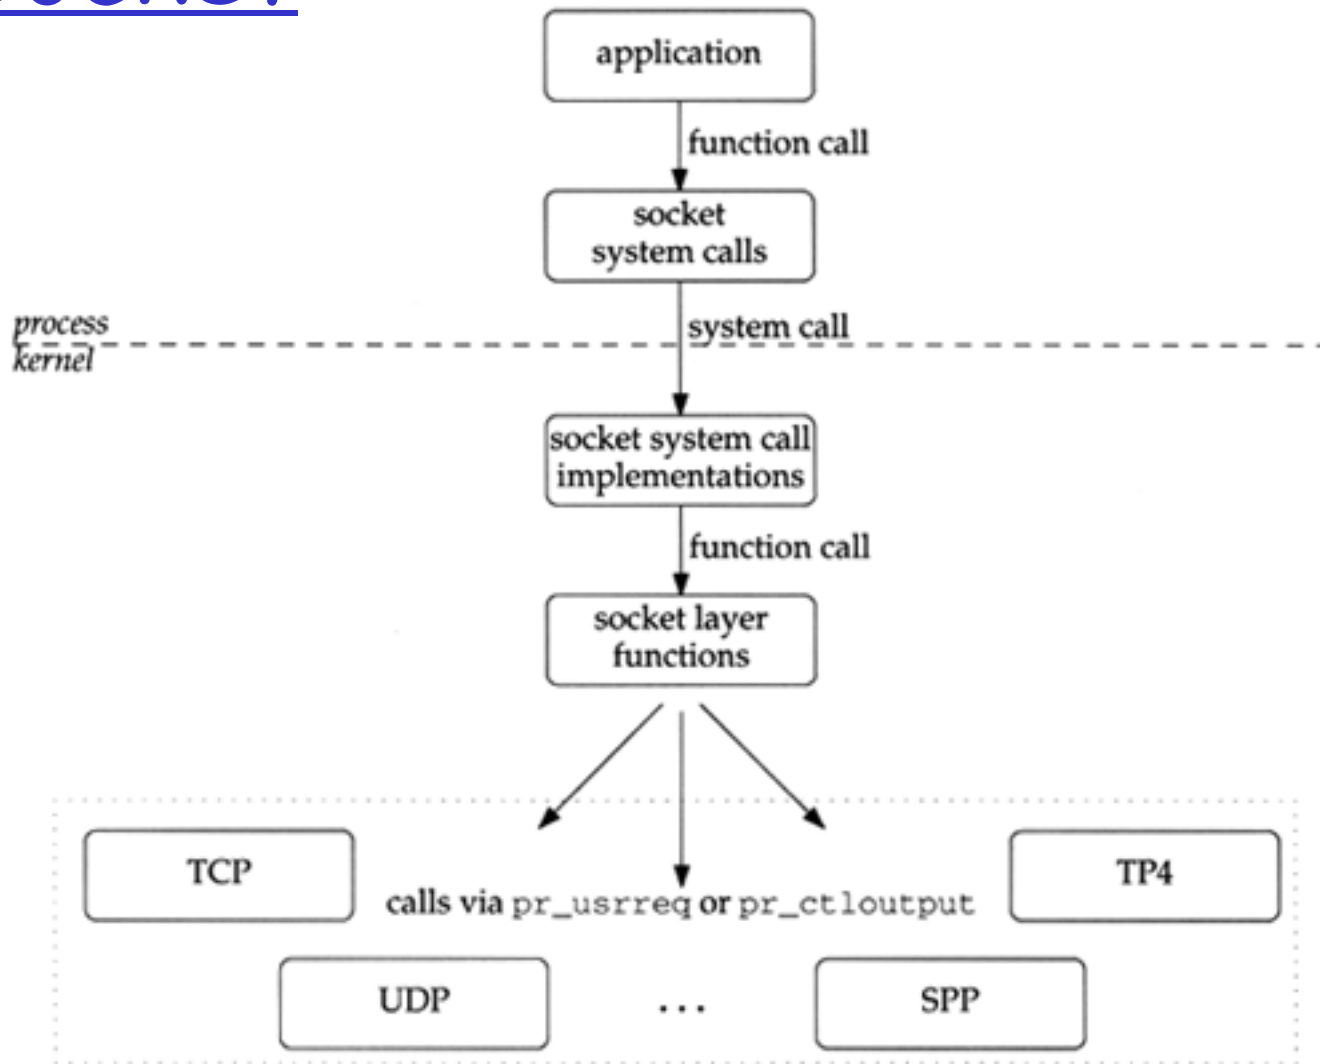
## □ Uses existing I/O programming interface as much as possible.



# Socket

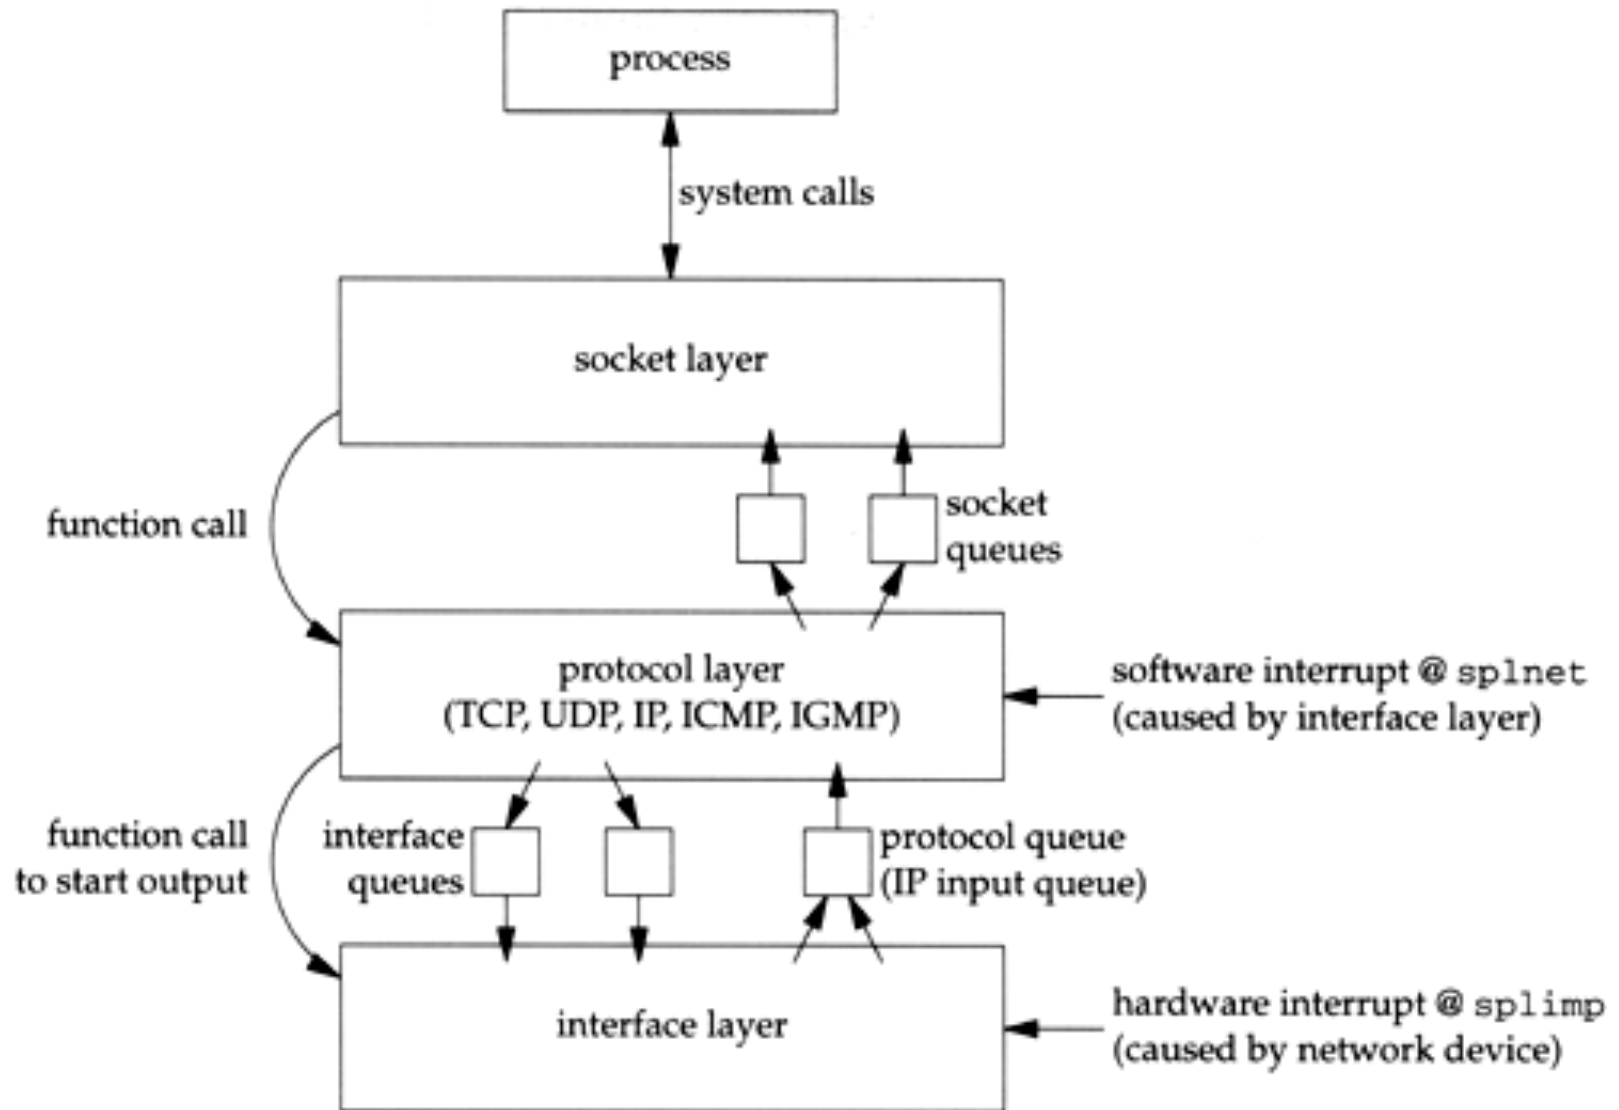
- ❑ A socket is an abstract representation of a communication endpoint.
- ❑ Sockets work with Unix I/O services just like files, pipes & FIFOs.
- ❑ Sockets (obviously) have special needs:
  - ❖ Establishing a connection
  - ❖ Specifying communication endpoint addresses

# Socket

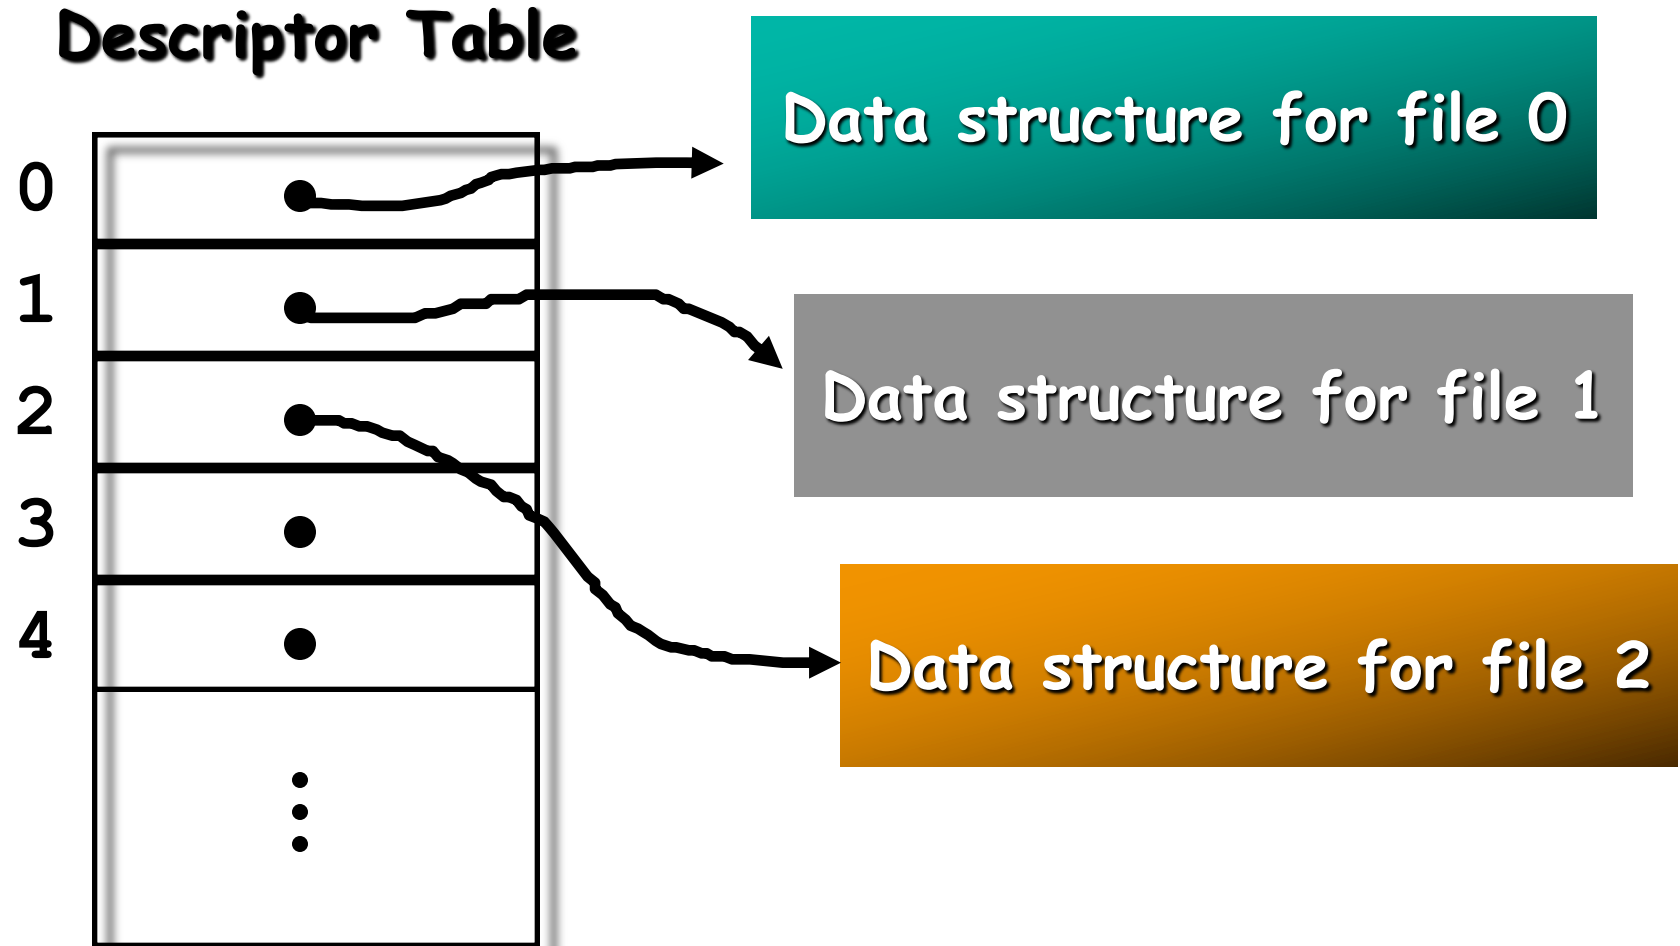




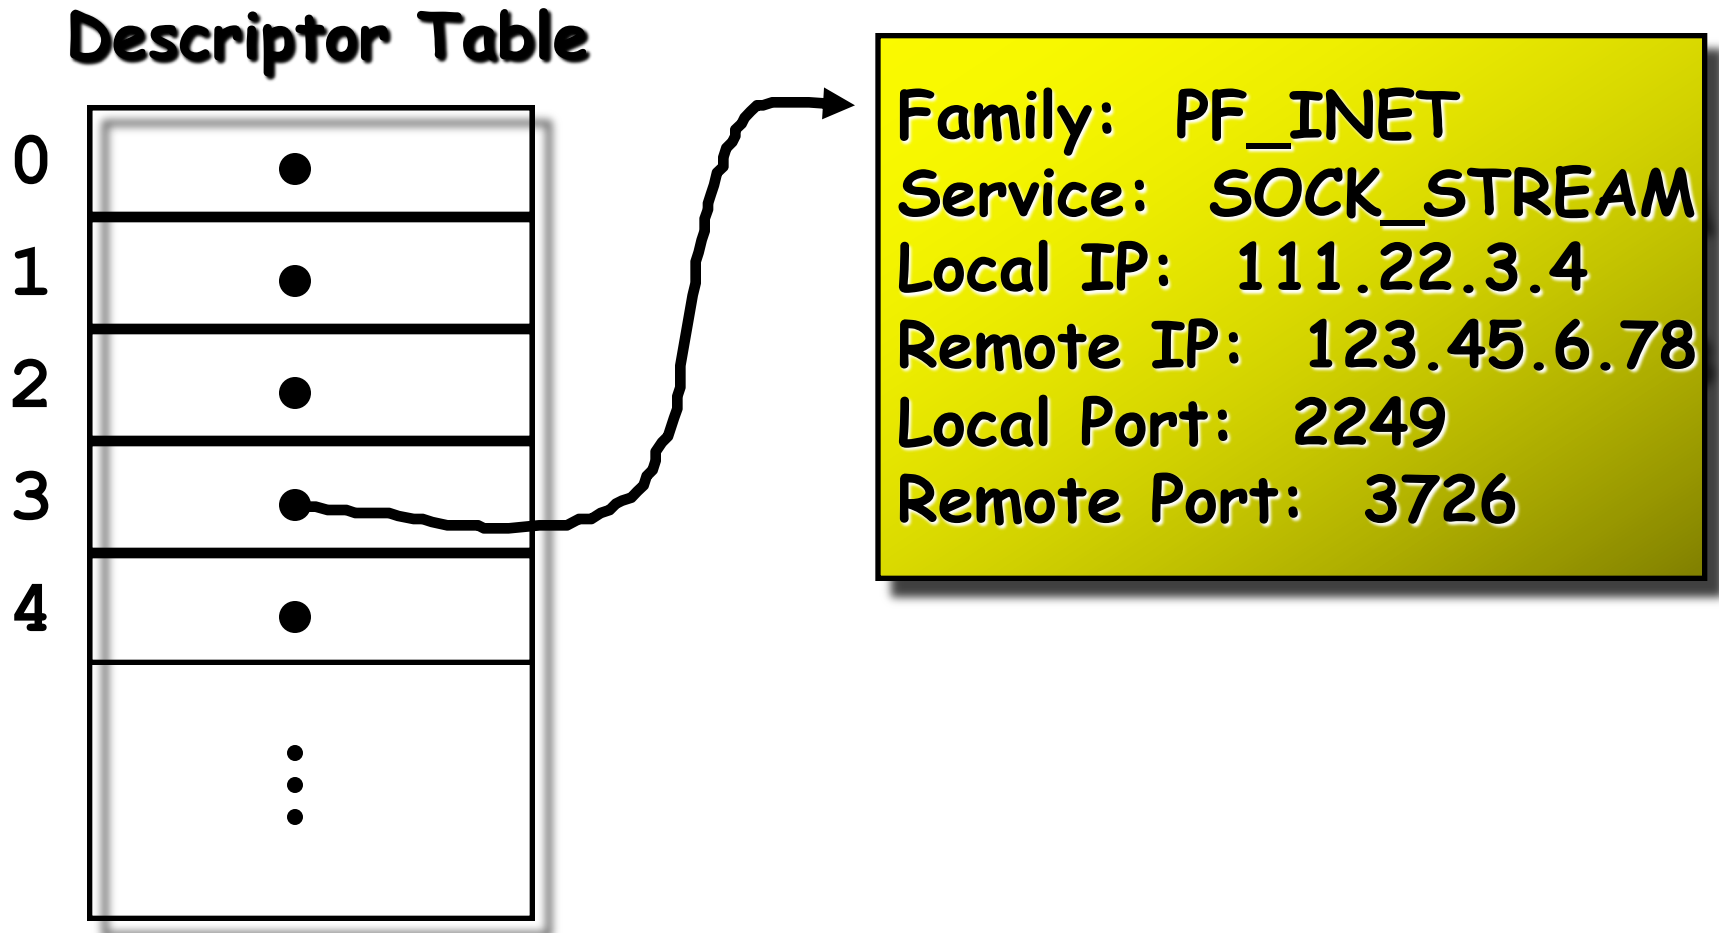
# Socket



# Unix Descriptor Table



# Socket Descriptor Data Structure



# Creating a Socket

```
int socket(int family, int type, int proto);
```

- ❑ **family** specifies the protocol family (PF\_INET for TCP/IP).
  - ❖ **AF\_INET** vs. **PF\_INET**
- ❑ **type** specifies the type of service (SOCK\_STREAM, SOCK\_DGRAM).
- ❑ **protocol** specifies the specific protocol (usually 0, which means *the default*).

## socket()

- ❑ **socket()** returns a socket descriptor (small integer) or -1 on error.
- ❑ **socket()** allocates resources needed for a communication endpoint
  - ❖ but it does not deal with endpoint addressing.

# Specifying an Endpoint Address

- ❑ Sockets API is generic.
- ❑ There must be a generic way to specify endpoint addresses.
- ❑ TCP/IP requires an IP address and a port number for each endpoint address.
- ❑ Other protocol suites (families) may use other schemes.

# Necessary Background Information: POSIX data types

**POSIX (Portable Operating System Interface of UNIX)**

<code>int8_t</code>	signed 8bit int
<code>uint8_t</code>	unsigned 8 bit int
<code>int16_t</code>	signed 16 bit int
<code>uint16_t</code>	unsigned 16 bit int
<code>int32_t</code>	signed 32 bit int
<code>uint32_t</code>	unsigned 32 bit int

`u_char, u_short, u_int, u_long`

Obsolete

# More POSIX data types

`sa_family_t`

address family

`socklen_t`

length of struct

`in_addr_t`

IPv4 address

`in_port_t`


IP port number



# Generic socket addresses

```
struct sockaddr {  
    uint8_t      sa_len;  
    sa_family_t  sa_family;  
    /* address family: AF_*** */  
    char          sa_data[14];  
    /* protocol-specific address */  
};
```

Used by kernel



- ❑ **sa\_family** specifies the address type.
- ❑ **sa\_data** specifies the address value.

# sockaddr

- ❑ An address that will allow me to use sockets to communicate with my family.

- ❑ address type **AF\_FAMILY**

- ❑ address values:

Daughter	1
----------	---

Wife	2
------	---

Mom	3
-----	---

Dad	4
-----	---

Sister	5
--------	---

Brother	6
---------	---

## AF\_FAMILY

- Initializing a sockaddr structure to point to Daughter:

```
struct sockaddr mary;
```

```
mary.sa_family = AF_FAMILY;
```

```
mary.sa_data[0] = 1;
```

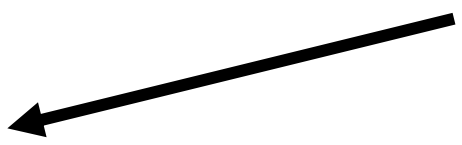
# AF\_INET

□ For AF\_FAMILY we only needed 1 byte to specify the address.

□ For AF\_INET we need:

- ❖ 16 bit port number
- ❖ 32 bit IP address

**IPv4 only!**



# struct sockaddr\_in (IPv4)

```
struct sockaddr_in {
```

```
    uint8_t            sin_len,
```

Length of  
structure (16)

```
    sa_family_t        sin_family;
```

AF\_INET

```
    in_port_t          sin_port;
```

16 bit  
Port number

```
    struct in_addr      sin_addr;
```

```
    char               sin_zero[8];
```

Make  
structure 16  
bytes

```
};
```

```
struct in_addr {
```

```
    in_addr_t          s_addr;
```

32 bit  
IPv4 address

```
};
```

# struct sockaddr\_in (IPv6)

```
struct sockaddr_in6 {  
    uint8_t            sin6_len;  
    sa_family_t        sin6_family;  
    in_port_t          sin6_port;  
    uint32_t           sin6_flowinfo;  
    struct in6_addr     sin6_addr;  
    uint32_t           sin6_scope_id;  
};  
  
struct in6_addr {  
    uint8_t            s6_addr[16];  
};
```

Length of structure (28)

AF\_INET6

Port number

Flow label

Scope of address

128 bit IPv6 address

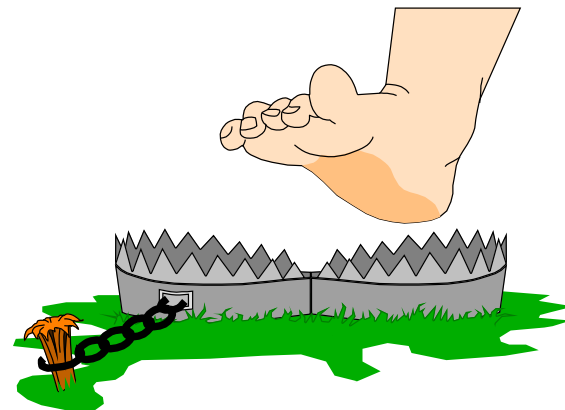
# Network Byte Order

□ All values stored in a `sockaddr_in` must be in network byte order.

❖ `sin_port` a TCP/IP port number.

❖ `sin_addr` an IP address.

!!! Common Mistake !!!



# Network Byte Order Functions

'h' : host byte order

's' : short (16bit)

'n' : network byte order

'l' : long (32bit)

```
uint16_t htons(uint16_t);
```

```
uint16_t ntohs(uint16_t);
```

```
uint32_t htonl(uint32_t);
```

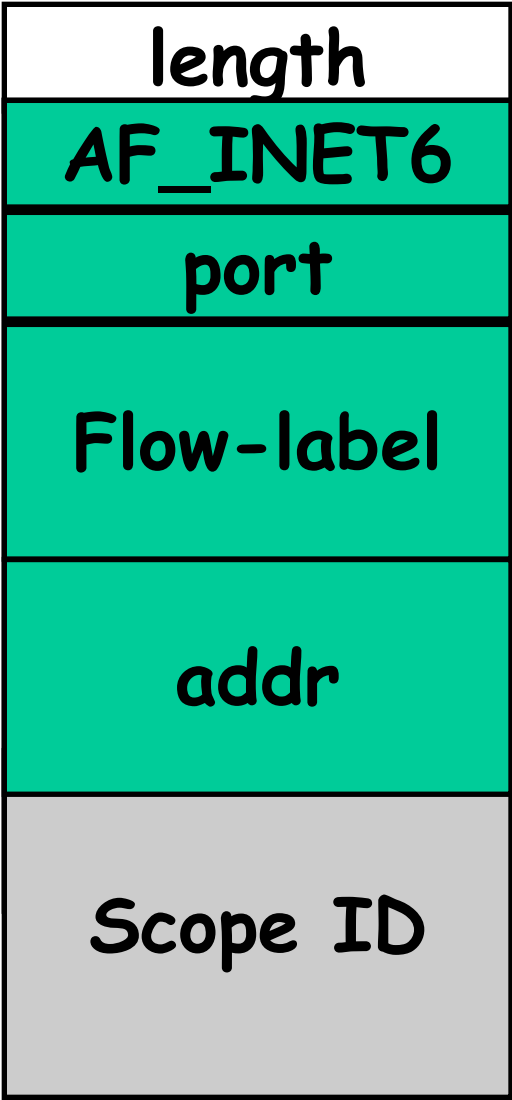
```
uint32_t ntohl(uint32_t);
```



# TCP/IP Addresses

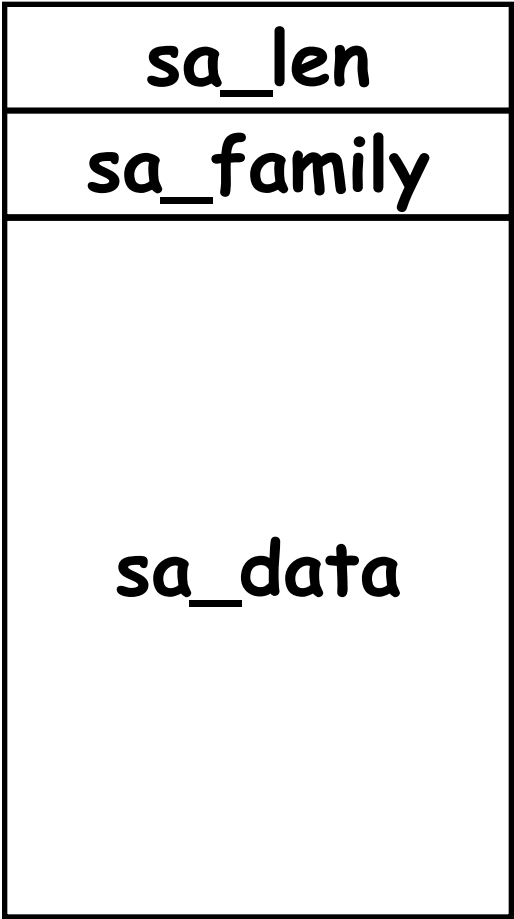
- ❑ We don't need to deal with **sockaddr** structures since we will only deal with a real protocol family.
- ❑ We can use **sockaddr\_in** structures.
- ❑ BUT: The C functions that make up the sockets API expect structures of type **sockaddr**.

# sockaddr\_in6



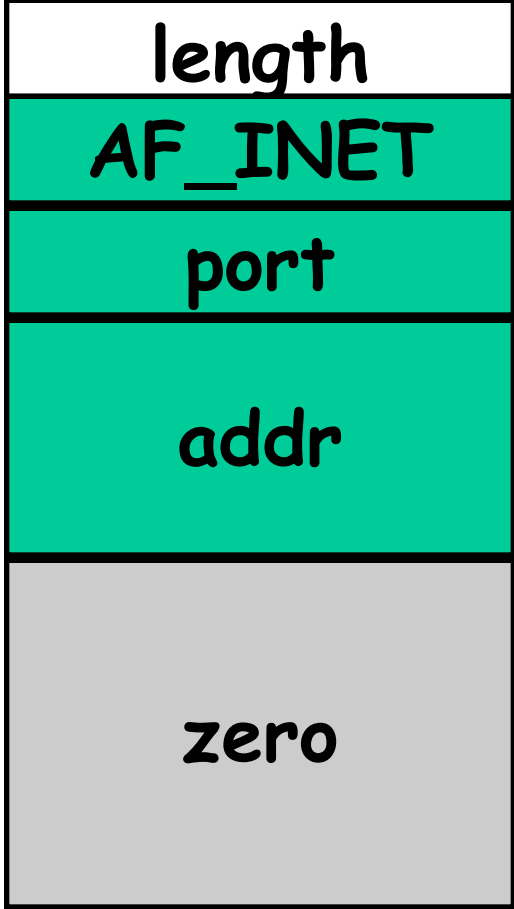
28 bytes

# sockaddr



variable

# sockaddr\_in



16 bytes

# Assigning An Address to A Socket

- The **bind()** is used to assign an address to an existing socket.

```
int bind( int sockfd,  
const!  → const struct sockaddr *myaddr,  
         int addrlen) ;
```

- **bind** returns 0 if successful or -1 on error.

## bind()

- Calling **bind()** assigns the address specified by the `sockaddr` structure to the socket descriptor.

- You can give **bind()** a `sockaddr_in` structure:

```
bind( mysock,  
      (struct sockaddr*) &myaddr,  
      sizeof(myaddr) );
```

# bind() Example

```
int mysock, err;  
struct sockaddr_in myaddr;  
  
mysock = socket(PF_INET, SOCK_STREAM, 0);  
myaddr.sin_family = AF_INET;  
myaddr.sin_port = htons( portnum );  
myaddr.sin_addr = htonl( ipaddress );  
  
err=bind(mysock, (sockaddr *) &myaddr,  
         sizeof(myaddr));
```

# Uses for bind()

- ❑ There are a number of uses for `bind()`:
  - ❖ Server would like to bind to a well known address (port number).
  - ❖ Client can bind to a specific port.
  - ❖ Client can ask the O.S. to assign *any available* port number.
- ❑ Clients typically don't care what port they are assigned.
- ❑ When you call `bind`, you can tell it to assign you any available port

# What Is My IP Address

- ❑ How can you find out what your IP address is so you can tell `bind()` ?
- ❑ There is no realistic way for you to know the right IP address to give `bind()`
  - ❖ what if the computer has multiple network interfaces?
- ❑ specify the IP address as: **INADDR\_ANY**, this tells the OS to take care of things.

# IPv4-only Address Conversion

`int inet_aton( char *, struct in_addr *);`

- ❖ Convert **ASCII** dotted-decimal IP address to **Network** byte ordered 32(IPv4) bit value.
- ❖ Returns 1 on success, 0 on failure.

`char *inet_ntoa(struct in_addr);`

- ❖ Convert **Network** byte ordered value to **ASCII** dotted-decimal (a string).



# IPv4&IPv6 Address Conversion

**int inet\_pton(int, const char \*, void \*);**

- ❖ (address\_family, string\_ptr, address\_ptr)
- ❖ Convert IP address string **presentation** to **Network byte ordered 32 (IPv4) or 128(IPv6) bit value**.
- ❖ Returns 1 on success, -1 on failure, 0 on invalid input

**char \*inet\_ntop(int, const void\*, char\*, size\_t);**

- ❖ (address\_family, address\_ptr, string\_ptr, length)
- ❖ Convert **Network byte ordered value** to IP address string **presentation**.
  - x:x:x:x:x:x:x:x or x:x:x:x:x:x:a.b.c.d

# Other Socket Functions

## □ General Use

- `read()`
- `write()`
- `close()`

## • Connection-oriented (TCP)

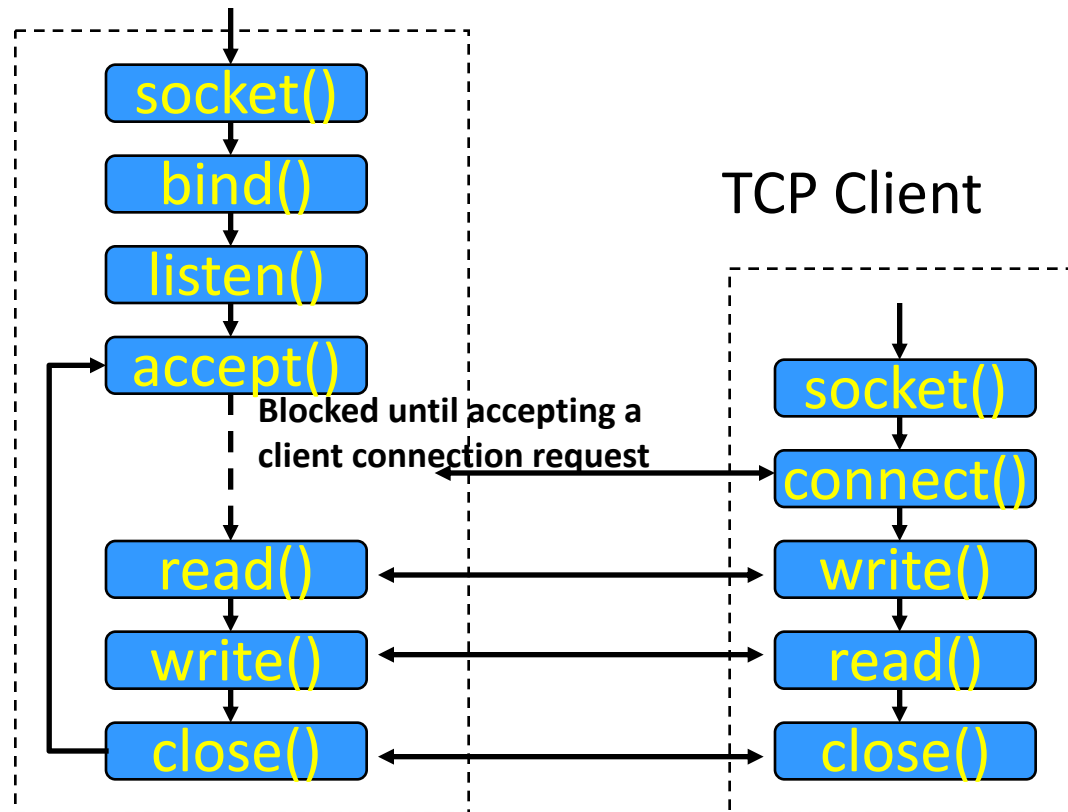
- `connect()`
- `listen()`
- `accept()`

## • Connectionless (UDP)

- `send()`
- `recv()`

# Basic TCP Socket Programming

TCP Server



# listen()

```
#include <sys/socket.h>
int listen(int s, int backlog);
```

## □ Parameters

### ❖ [in] socket

- Specifies the file descriptor associated with the socket.

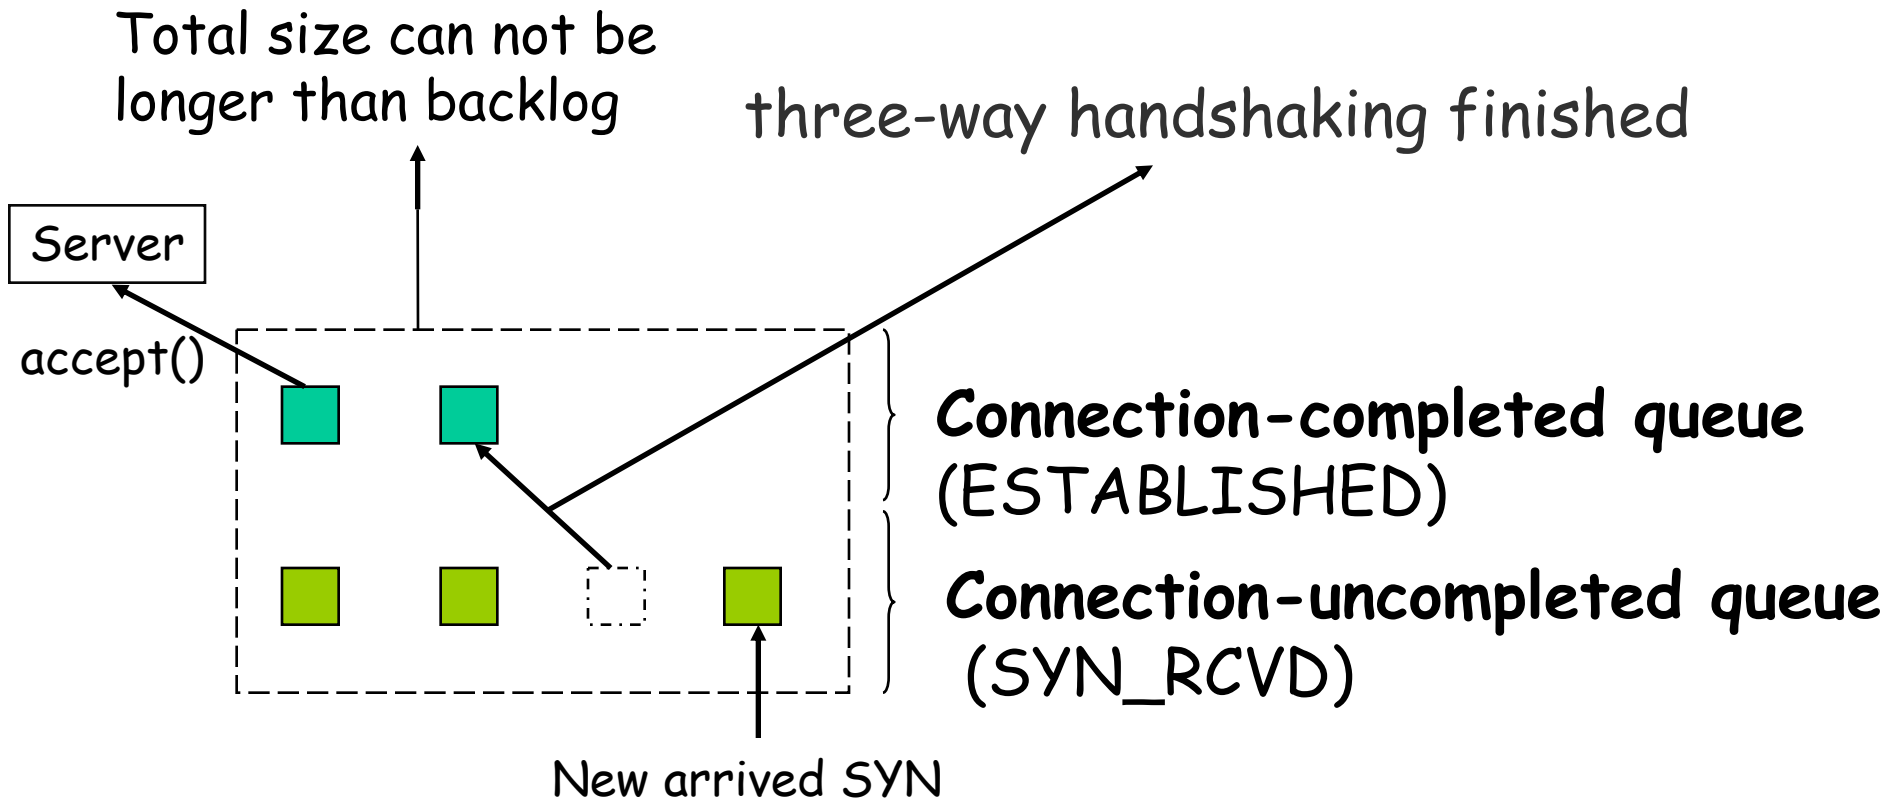
### ❖ [in] backlog

- To limit the number of outstanding connections in the socket's listen queue. Implementations shall support values of backlog up to `SOMAXCONN`, defined in `<sys/socket.h>`.

## □ Returns

- ❖ Upon successful completions, **listen()** shall return 0; otherwise, -1 shall be returned and `errno` set to indicate the error.

# listen()-backlog



**Two Queues maintained by TCP**

# listen()-backlog

## □ Linux IPv4 Implementation

- ❖ The size of **connection-uncompleted queue** is defined by **backlog**. But the value of backlog cannot be larger than the value stored in **/proc/sys/net/core/somaxconn** with default value 128.
- ❖ The size of **connection-uncompleted queue** is defined by **/proc/sys/net/ipv4/tcp\_max\_syn\_backlog**

# accept()

```
#include <sys/types.h>
#include <sys/socket.h>
int accept(int sockfd, struct sockaddr
*addr, socklen_t *addrlen);
```

❑ Accept an incoming connection on a listening socket

## ❑ Parameters

❖ [in] sockfd

- The listen()ing socket descriptor.

❖ [out] addr

- Filled in with the address of the site that's connecting to you.

❖ [out] addrlen

- Filled in with the sizeof() the structure returned in the addr parameter.

## ❑ Returns

- ❖ The newly connected socket descriptor, or -1 on error, with errno set appropriately.

## connect()

```
#include <sys/types.h>
#include <sys/socket.h>
int connect(int sockfd, const struct
sockaddr *serv_addr, socklen_t addrlen);
```

❑ Once building a socket descriptor with **socket()**, you can **connect()** that socket to a remote server.

### ❑ Parameters

#### ❖ [in] sockfd

- The socket descriptor to

#### ❖ [in] addr

- Filled in with the address of the site you want to connect to.

#### ❖ [in] addrlen

- Specifies the length of the sockaddr structure pointed to by the address argument..

### ❑ Returns

- ❖ 0 if succeed or -1 on error, with errno set appropriately



## send()

```
#include <sys/types.h>
#include <sys/socket.h>
ssize_t send(int sockfd, const void *buf, size_t
len, int flags);
```

❑ Send a message on a socket, used for TCP  
SOCK\_STREAM connected sockets

### ❑ Parameters

- ❖ [in] sockfd: Specifies the socket file descriptor.
- ❖ [in] buf: Points to the buffer containing the message to send.
- ❖ [in] len: Specifies the length of the message in bytes.
- ❖ [in] flags: Specifies the type of message reception. Support for values other than 0 is not implemented yet.

### ❑ Returns

- ❖ The number of bytes sent if succeed or -1 on error, with errno set appropriately.

# sendto()

```
#include <sys/types.h>
#include <sys/socket.h>
ssize_t sendto(int sockfd, const void *buf,
size_t len, int flags, const struct sockaddr
*to, socklen_t tolen);
```

- ❑ Send a message on a socket, used for UDP SOCK\_DGRAM unconnected datagram sockets.
- ❑ Parameters
  - ❖ [in] sockfd: Specifies the socket file descriptor.
  - ❖ [in] buf: Points to the buffer containing the message to send.
  - ❖ [in] len: Specifies the length of the message in bytes.
  - ❖ [in] flags: Specifies the type of message reception. Support for values other than 0 is not implemented yet.
  - ❖ [in] address: Points to a sockaddr structure containing the destination address. The length and format of the address depend on the address family of the socket.
  - ❖ [in] address\_len: Specifies the length of the sockaddr structure

# Socket核心函数-recv()

- 功能:
  - ① 从TCP接收数据,返回实际接收数据长度, 出错返回-1。
  - ② 如果没有数据将阻塞, 如果收到的数据大于缓存大小, 多余数据将丢弃。
- 参数:
  - ① Sockfd:套接字描述符
  - ② Buf:指向内存块的指针
  - ③ Buf\_len:内存块大小, 以字节为单位
  - ④ flags:一般为0

```
int recv(int sockfd, void *buf, int buf_len,unsigned  
int flags);
```

# Socket核心函数-recv()

**int recv(int sockfd, void \*buf, int buf\_len,unsigned  
int flags);**

**e.g. recv(sockfd, buf, 8192, 0)**

# Socket核心函数-recvfrom()

- 功能:
  - ① 从UDP接收数据，返回实际接收的字节数，失败返回-1
- 参数:
  - ① sockfd:套接字描述符
  - ② buf:指向内存块的指针
  - ③ buf\_len:内存块大小，以字节为单位
  - ④ flags:一般为0
  - ⑤ from:远端的地址，IP地址和端口号
  - ⑥ fromlen:远端地址长度

**int recvfrom(int sockfd, void \*buf, int buf\_len,  
unsigned int flags, struct sockaddr \*from, int fromlen)**

# Socket核心函数- recvfrom()

**int recvfrom(int sockfd, void \*buf, int buf\_len,  
unsigned int flags, struct sockaddr \*from, int  
fromlen)**

**e.g. recvfrom(sockfd, buf, 8192, 0, (struct  
sockaddr \*)&address, sizeof(address));**

# Socket核心函数-close()

## ❑ 功能:

- ① 关闭**socket**，实际上是关闭文件描述符
- ② 如果只有一个进程使用，立即终止连接并撤销该套接字，如果多个进程共享该套接字，将引用数减一，如果引用数降到零，则撤销它。

## ❑ 参数:

- ① **Sockfd**:套接字描述符

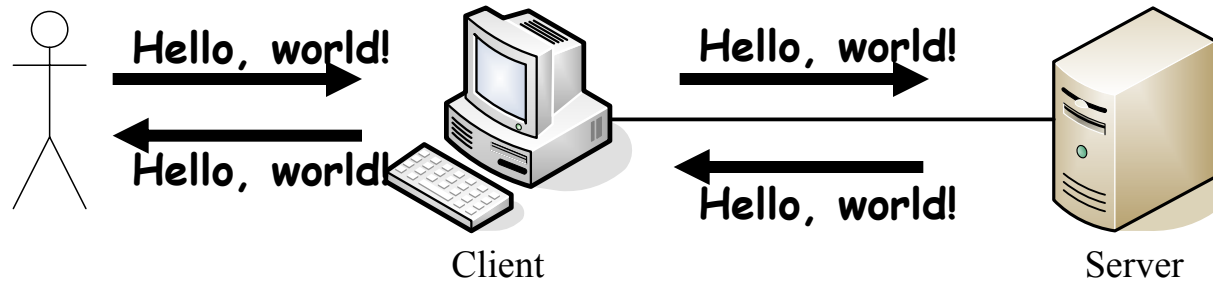
```
#include <unistd.h>

int close (int sockfd);
```

Returns: 0 if OK, -1 on error

# Iterative Server

## □ Echo process





# Server

```
1 #include      "unp.h"

2 int
3 main(int argc, char **argv)
4 {
5     int      listenfd, connfd;
6     pid_t    childpid;
7     socklen_t clilen;
8     struct sockaddr_in cliaddr, servaddr;

9     listenfd = Socket (AF_INET, SOCK_STREAM, 0);

10    bzero(&servaddr, sizeof(servaddr));
11    servaddr.sin_family = AF_INET;
12    servaddr.sin_addr.s_addr = htonl (INADDR_ANY);
13    servaddr.sin_port = htons (SERV_PORT);
```

# Server

```
14      Bind(listenfd, (SA *) &servaddr, sizeof(servaddr));

15      Listen(listenfd, LISTENQ);

16      for ( ; ; ) {
17          clilen = sizeof(cliaddr);
18          connfd = Accept(listenfd, (SA *) &cliaddr, &clilen);
19          str_echo(connfd);    /* process the request */
20          Close(connfd);
21      }
22 }
```

# Server

```
1 #include    "unp.h"

2 void
3 str_echo(int sockfd)
4 {
5     ssize_t n;
6     char    buf[MAXLINE];

7     again:
8     while ( (n = read(sockfd, buf, MAXLINE)) > 0)
9         Writen(sockfd, buf, n);

10    if (n < 0 && errno == EINTR)
11        goto again;
12    else if (n < 0)
13        err_sys("str_echo: read error");
14 }
```

# Client

```
1 #include    "unp.h"

2 int
3 main(int argc, char **argv)
4 {
5     int      sockfd;
6     struct sockaddr_in servaddr;

7     if (argc != 2)
8         err_quit("usage: tcpcli <IPaddress>");

9     sockfd = Socket(AF_INET, SOCK_STREAM, 0);

10    bzero(&servaddr, sizeof(servaddr));
11    servaddr.sin_family = AF_INET;
12    servaddr.sin_port = htons(SERV_PORT);
13    Inet_pton(AF_INET, argv[1], &servaddr.sin_addr);
```

# Client

```
14  Connect(sockfd, (SA *) &servaddr, sizeof(servaddr));  
  
15      str_cli(stdin, sockfd);      /* do it all */  
  
16      exit(0);  
17 }
```

# Client

```
2 void
3 str_cli(FILE *fp, int sockfd)
4 {
5     char    sendline[MAXLINE], recvline[MAXLINE];

6     while (Fgets(sendline, MAXLINE, fp) != NULL) {

7         Writen(sockfd, sendline, strlen (sendline));

8         if (Readline(sockfd, recvline, MAXLINE) == 0)
9             err_quit("str_cli: server terminated
premadurely");

10         Fputs(recvline, stdout);
11     }
12 }
```