

Network Computing & Programming

Lecture 4

Key Problems of

Client & Server Design

2019-2020-1

Slides are modified from Xiang(Shaun) Zhang

Issues in C/S Programming

- ① Identifying the Server.
- ② Looking up an IP address.
- ③ Looking up a well known port name.
- ④ Specifying a local IP address.
- ⑤ UDP/TCP client design.

Identifying the Server

□ Options:

- ① Hard-coded into the client program.
- ② Require that the user identify the server.
- ③ Read from a configuration file.
- ④ Use a separate protocol/network service to lookup the identity of the server.

Identifying a TCP/IP server

❑ Need an **IP address**, **protocol** and **port**.

❖ Usually use host names instead of IP addresses

- Two parameter: **ostec.uestc.edu.cn http**
- One parameter: **ostec.uestc.edu.cn:http**

Converting an IP address

```
in_addr_t inet_addr(const char *strptr);
```

- ❖ Convert **ASCII** dotted-decimal IP address to **network** byte ordered 32-bit (IPv4) value. Just like **inet_aton()** learned before.
- ❖ Returns a 32-bit binary network byte ordered IPv4 address and **INADDR_NONE** on error

(...)

```
struct sockaddr_in dest;
```

```
memset(&dest, '\0', sizeof(dest));
```

```
dest.sin_addr.s_addr = inet_addr("202.115.16.62");
```

(...)

Looking up an IP address

**struct hostent *gethostbyname(const char
*hostname); (but obsolete, use getaddrinfo() now)**

```
struct hostent{  
    Char    *h_name;  
    Char    ** h_aliases;  
    Short   h_addrtype;  
    Short   h_length;  
    Char    ** h_addr_list;};
```

(...)

struct hostent *remoteHost;

remoteHost = gethostbyname(hostname)

(...)

Looking up a well known port name

```
struct servent *getservbyname(const char * name,  
    const char *proto);
```

```
struct servent{  
    char    *s_name;  
    char    **s_aliases;  
    int     s_port;  
    char    *s_proto;};
```

(...)

```
struct servent *remoteSvr;
```

```
remoteSvr = getservbyname("domain", "udp")
```

(...)

Looking up a protocol

struct protoent *getprotobyname(int proto)

❖ Look up /etc/protocols to get result.

```
struct protoent {  
    char    *p_name;  
    char    *p_aliases;  
    short   *p_proto;};
```

(...)


struct protoent *pptr;

pptr = getprotobyname("udp")

(...)

Specifying a Local Address

- ❑ When a client creates and binds a socket, it must specify **a local port** and **IP address**
- ❑ Typically a client doesn't care what port it is working on:

 `haddr->port = htons(0);`
Give me any available port !

Local IP address

- A client can also ask the operating system to take care of specifying the local IP address:

```
haddr->sin_addr.s_addr=  
    htonl(INADDR_ANY);
```



Give me the appropriate address

UDP Client Design

- ❑ Establish server address (IP and port).
- ❑ Allocate a socket.
- ❑ Specify that any valid local port and IP address can be used.
- ❑ Communicate with server (send, recv)
- ❑ Close the socket.

Connected mode UDP

- ❑ A UDP client can call `connect()` to establish the address of the server
- ❑ The UDP client can then use `read()` and `write()` or `send()` and `recv()`
- ❑ A UDP client using a connected mode socket can only talk to one server
 - ❖ using the connected-mode socket

TCP Client Design

- ❑ Establish server address (IP and port).
- ❑ Allocate a socket.
- ❑ Specify that any valid local port and IP address can be used.
- ❑ Call connect()
- ❑ Communicate with server (read, write).
- ❑ Close the connection.

Closing a TCP socket

- ❑ Many TCP based application protocols support
 - ❖ multiple requests and/or
 - ❖ variable length requests over a single TCP connection.

- ❑ How does the server know when the client is done ?
 - ❖ and it is OK to close the socket ?

Partial Close

- ❑ One solution is for the client to shut down only it's writing end of the socket.
- ❑ `Shutdown()` provides this function.
`shutdown(int s, int direction);`
 - ❖ direction can be 0 to close the reading end or 1 to close the writing end.
 - ❖ shutdown sends info to the other process!

TCP sockets programming

□ Common problem areas:

- ❖ null termination of strings.
- ❖ reads don't correspond to writes.
- ❖ synchronization (including close()).
- ❖ ambiguous protocol.

TCP Reads

- ❑ Each call to `read()` on a TCP socket returns any available data
 - ❖ up to a maximum
- ❑ TCP buffers data at both ends of the connection.
- ❑ *You must be prepared to accept data 1 byte at a time from a TCP socket!*

Server Design

Iterative
Connectionless

Iterative
Connection-Oriented

Concurrent
Connectionless

Concurrent
Connection-Oriented

Concurrent vs. Iterative

Concurrent

Large or variable size requests
Harder to program
Typically uses more system resources

Iterative

Small, fixed size requests
Easy to program

Connectionless vs. Connection-Oriented

Connection-Oriented

EASY TO PROGRAM

transport protocol handles the tough stuff.
requires separate socket for each connection.

Connectionless

less overhead
no limitation on number of clients

Statelessness

- ❑ ***State***: Information that a server maintains about the status of ongoing client interactions.
- ❑ Connectionless servers that keep state information must be designed carefully!

Messages can be duplicated!

The Dangers of Statelessness

- ❑ Clients can go down at any time.
- ❑ Client hosts can reboot many times.
- ❑ The network can lose messages.
- ❑ The network can duplicate messages.

Concurrent Server Design Alternatives

- ❑ One child per client
- ❑ Spawn one thread per client
- ❑ Preforking multiple processes
- ❑ Prethreaded Server

One child per client

❑ Traditional Unix server:

- ❖ TCP: after call to `accept()` , call `fork()` .
- ❖ UDP: after `recvfrom()` , call `fork()` .
- ❖ Each process needs only a few sockets.
- ❖ Small requests can be serviced in a small amount of time.

❑ Parent process needs to clean up after children!!!!

- ❖ call `wait()`

One thread per client

- ❑ Almost like using fork
 - ❖ call `pthread_create` instead
- ❑ Using threads makes it easier to have sibling processes share information
 - ❖ less overhead
- ❑ Sharing information must be done carefully
 - ❖ use `pthread_mutex`

Prefork () 'd TCP Server

- ❑ Initial process creates socket and binds to well known address.
- ❑ Process now calls `fork ()` a bunch of times.
- ❑ All children call `accept ()` .
- ❑ The next incoming connection will be handed to one child.

Preforking

- ❑ Having too many preforked children can be bad.
- ❑ Using dynamic process allocation instead of a hard-coded number of children can avoid problems.
- ❑ Parent process just manages the children
 - ❖ doesn't worry about clients

Sockets library vs. system call

- ❑ A preforked TCP server won't usually work the way we want if *sockets* is not part of the kernel:
 - ❖ calling `accept()` is a library call, not an atomic operation.
- ❑ We can get around this by making sure only one child calls `accept()` at a time using some locking scheme.

Prethreaded Server

- ❑ Same benefits as preforking.
- ❑ Can also have the main thread do all the calls to `accept()`
 - ❖ and hand off each client to an existing thread

What's the best server design for my application?

□ Many factors:

- ❖ expected number of simultaneous clients
- ❖ Transaction size
 - time to compute or lookup the answer
- ❖ Variability in transaction size
- ❖ Available system resources
 - perhaps what resources can be required in order to run the service

Server Design

- ❑ It is important to understand the issues and options.
- ❑ Knowledge of queuing theory can be a big help.
- ❑ You might need to test a few alternatives to determine the best design.